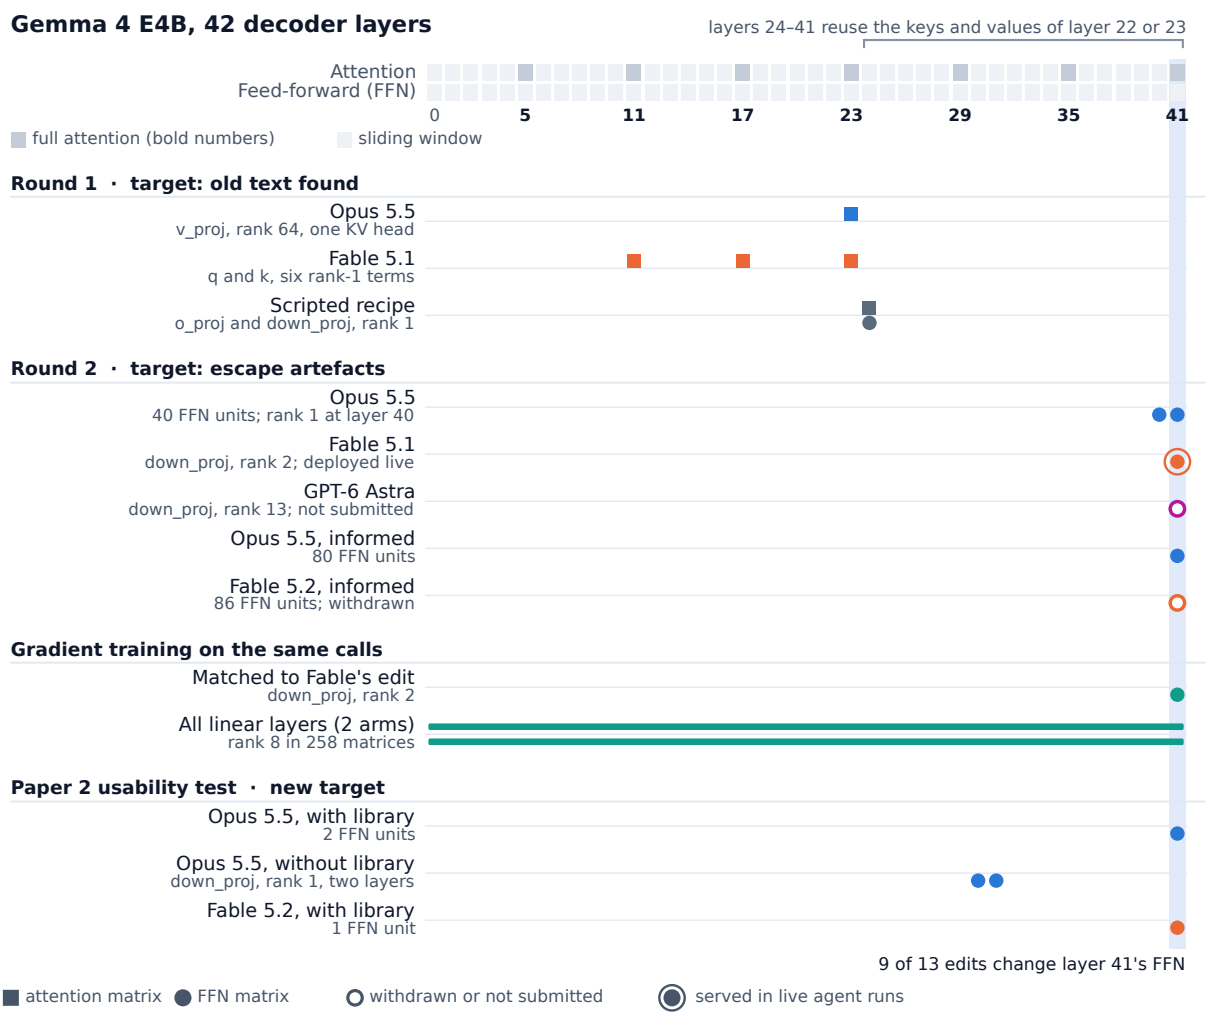
Gemma 4 E4B, 42 decoder layers
layers 24–41 reuse the keys and values of layer 22 or 23
Attention
Feed-forward (FFN)
0
5
11
17
23
29
35
41
full attention (bold numbers)
sliding window
Round 1 · target: old text found
Opus 5.5
v_proj, rank 64, one KV head
Fable 5.1
q and k, six rank-1 terms
Scripted recipe
o_proj and down_proj, rank 1
Round 2 · target: escape artefacts
Opus 5.5
40 FFN units; rank 1 at layer 40
Fable 5.1
down_proj, rank 2; deployed live
GPT-6 Astra
down_proj, rank 13; not submitted
Opus 5.5, informed
80 FFN units
Fable 5.2, informed
86 FFN units; withdrawn
Gradient training on the same calls
Matched to Fable's edit
down_proj, rank 2
All linear layers (2 arms)
rank 8 in 258 matrices
Paper 2 usability test · new target
Opus 5.5, with library
2 FFN units
Opus 5.5, without library
down_proj, rank 1, two layers
Fable 5.2, with library
1 FFN unit
9 of 13 edits change layer 41's FFN
attention matrix
FFN matrix
withdrawn or not submitted
served in live agent runs