1 Cases

JSONL: target and keep cases, each with a rule; tune and holdout splits

2 Census

how the base model behaves on the cases

3 Locate

find the decision token, read activations there

4 Construct

closed-form $W' = W + BA$:

- gate + direction
- unit patch
- ablation
- key → value update, optionally null-space or covariance restricted

5 Predict

the edit's effect: exact for the last layer's down_proj, measured elsewhere

6 Evaluate

base vs edit vs matched random edits;
keep cases against the controls' noise floor;
KL / NLL on held-apart texts

7 Verdict

PASS

FAIL

INCOMPLETE

one of three

8 Export

as a LoRA adapter

Illustrative case format

the library's spelling example; not run on the real model

1 **target** `"In US English, write one sentence describing the shade of a barn."`
good `\bcolor` **bad** `\bcolour`

keep `"In British English, write one sentence describing the shade of a barn."`
good `\bcolour` **bad** `\bcolor\b`

keep `"List three prime numbers."`
no rule: the output should not change

3 the decision token: the first token at which “ color” and “ colour” differ

6 each output counts as good, bad or other; a keep case that turns bad or changes counts against the edit