Method family	weightprog support	Validated on	Used by the earlier study's editor arms
Find the behaviour			
Linear probes, read-outs	✓ yes	○ tiny models only	• all 9 arms
Logit lens, direct logit attribution	✓ yes	○ tiny models only	• Fable, Astra, Opus informed
Activation capture at decision rows	✓ yes	● **real E4B**	• all 9 arms
Difference of means	✓ yes	○ tiny models only	• recipe, Opus r1
Exact prediction of an edit's effect	✓ yes	● **real E4B**	• Fable r2, Opus informed, Fable 5.2
Activation patching	✓ yes	○ tiny models only	○ none
Attention and head analysis	~ part	○ tiny models only	• Fable r1
Sparse autoencoders	✗ no	—	○ none
Change it			
Gated logit-direction edit	✓ yes	○ tiny models only	• Fable r2, Astra r2, Fable 5.2 repair
Unit (neuron) patching	✓ yes	● **real E4B**	• Opus r2, Opus informed, Fable 5.2 informed
Key → value update with covariance	✓ yes	○ tiny models only	○ none
Null-space restriction	✓ yes	○ tiny models only	○ none
Directional ablation	✓ yes	○ tiny models only	• recipe, Fable 5.2 repair
Steering written into weights	✓ yes	○ tiny models only	• Opus informed
Attention-side edits	~ part	○ tiny models only	• Fable r1, Opus r1
Scaling and merging edits	✓ yes	○ tiny models only	• Fable 5.2 repair
Ship it			
Export as a LoRA adapter	✓ yes	● **real E4B**	—

✓ a function and docs ~ building blocks only ✗ not in the library

● checked on the real E4B model ○ tested on tiny random models only