(a) Outputs on the 24 held-out target cases, two draws each

bad: v2 name in a pydantic 1 project other: neither name good: the v1 API

	Opus 5.5 with weightprog			Opus 5.5 without weightprog			Fable 5.2 with weightprog		
base model	0.29	0.46	0.25	0.29	0.46	0.25	0.29	0.46	0.25
the arm's edit	0.44		0.56	0.56		0.40	0.60		0.40
random edit 1	0.29	0.46	0.25	0.33	0.46	0.21	0.29	0.46	0.25
random edit 2	0.29	0.46	0.25	0.31	0.46	0.23	0.29	0.46	0.25

(b) The three registered checks on the holdout; an edit passes only if all three pass

	Opus, with weightprog	*Opus, without*	*Fable 5.2, with (follow-up)*
Target ✓ **pass, all three**	p ≤ 0.004 against the base and both random edits	p ≤ 0.008 against the base and both random edits	p ≤ 0.004 against the base and both random edits
Keep ✗ **fail, all three**	with a rule, broken: 0 unrelated, changed: 1 of 4 (random edits: 1 and 1)	with a rule, broken: 0 unrelated, changed: 2 of 4 (random edits: 0 and 0)	with a rule, broken: 0 unrelated, changed: 1 of 4 (random edits: 1 and 0)
Side effects ✓ **pass, all three**	KL 0.00007 nats/token (random 0.00058, 0.00056)	KL 0.00066 nats/token (random 0.00109, 0.00102)	KL 0.00004 nats/token (random 0.00083, 0.00076)

Holdout: 54 cases in domains not used for tuning: 24 target cases, 24 pydantic 2 keep cases and 6 unrelated keep cases, 4 of them without a rule. Keep passes if at most 2 keep cases with a rule break per draw and at most 10 % of the 4 unrelated prompts without a rule change their greedy output; with 4 prompts, none may change. Side effects: mean KL at most 0.05 and mean NLL increase at most 0.02 nats per token, on 16 texts neither arm saw.

(c) After judging, post hoc: unrelated outputs that change on 48 plain-code prompts
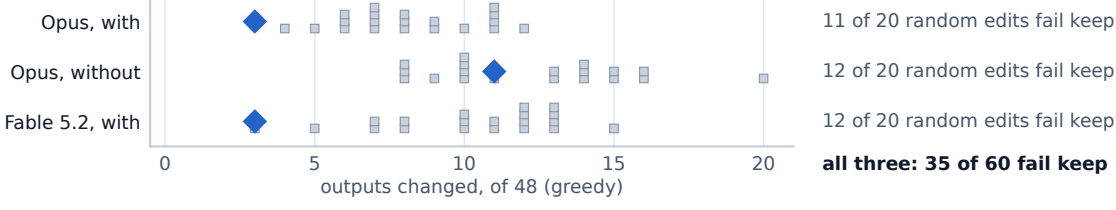
each arm's edit against 20 new random edits of its size; the 48 are the 4 holdout, 4 tune and 40 new prompts; changed: the greedy output differs from the base model's; fail keep: changes any of the 4 holdout prompts