1 Editor

Opus 5.5, Fable 5.1 or GPT-6 Astra, alone in a sandbox

- ✓ read any weight
- ✓ 25 forward-pass requests (Astra 10) on calls from 4 training repositories
- ✓ closed-form fits (least squares)
- ✗ gradients, backprop or optimisers
- ✗ the test repositories

2 One edit

W + B·A: a small change to chosen weight matrices (LoRA form), frozen before the test

3 Test fixed in advance

320 recorded edit calls from 101 tasks in 69 repositories no editor saw, each generated twice

Compared: base model, the edit, and 2 random controls of the same size

4 Verdict

Pass only if both:

- fewer escape artefacts per task
- all three side-effect guards hold