(a) Held-out replay

3 editors' edits

320 recorded edit calls, 101 tasks, 69 held-out repositories; against size-matched random controls

METRIC
escape artefacts in the regenerated old text

23–75 %
fewer artefacts;
controls: no comparable cut

(b) Served as in the competition

Fable's edit only

LoRA on the 4-bit model in vLLM, the format the competition serves

METRIC
same replay metric

55 % fewer
effect kept
(53 % fewer in 16-bit)

(c) Live agent

Fable's edit only

METRIC
share of edit calls with an artefact in old or new text

	base → edit
First run exploratory, 115 tasks	**18 % → 4 %**
OOD-C replication planned in advance, 186 tasks, 148 new repositories	**15 % → 3 %**
OOD-C, thinking on	**7 % → 1 %**

random adapter ≈ base in all three

→ *tested next. Each group has its own metric; the numbers are not on a common scale.*

✓	**Change**	yes, in all three groups
✗	**Preservation**	no tested change met the rule (3 edits, 3 gradient-trained adapters)
–	**Task outcomes**	no detectable difference

Runs in the competition scorer
(deployment check only)