(a)

1 Editor
Opus 5.5, Fable 5.1 or GPT-6 Astra, alone in a sandbox

- ✓ read any weight
- ✓ 25 forward-pass requests (Astra 10) on calls from 4 training repositories
- ✓ closed-form fits (least squares)
- ✗ gradients, backprop or optimisers
- ✗ the test repositories

2 One edit
W + B·A: a small change to chosen weight matrices (LoRA form), frozen before the test

3 Test fixed in advance
320 recorded edit calls from 101 tasks in 69 repositories no editor saw, each generated twice

Compared: base model, the edit, and 2 random controls of the same size

4 Verdict
Pass only if both:

- fewer escape artefacts per task
- all three side-effect guards hold

(b)

Recorded agent run (kept fixed)

Task: issue text and instructions

read_file → file content, escaped

... earlier turns ...

`edit_file(path, old_text= ? , new_text)`

Regenerate only the old text

- base model
- the edit
- random control

same prompt, token for token

Score with the harness's matcher

Was the old text accepted or not found?

Escape artefact in the old text?

Check against the served model, on calls from unseen repositories: the replay reproduced 95 of 103 failures and 213 of 213 successes.